




Getting start with Azure OpenAI

Sergio Zenatti Filho
Sr Cloud Solution Architect





Sergio Zenatti Filho

Senior Cloud Solution Architect
at Microsoft

Sergio has over 20 years of experience designing and delivering Data and Analytics Solutions. He has extensive experience in the Microsoft Data and Analytics Platform in the cloud and also on-premises. Sergio is passionate about learning new technology and helping customers to define the best solution for their business.

Connect



in /sergiozenatti

Agenda

- Azure OpenAI Overview
- Why Azure OpenAI
- Concepts
- Demos
- Use Cases
- Next Steps
- Q&A



*Ensure that artificial
general intelligence (AGI)
benefits humanity*



*Empower every person and
organization on the planet
to achieve more*

GPT-3.5 and GPT-4

Text

ChatGPT

Conversation

Codex

Code

DALL·E 2

Images

Azure AI

Applications



Partner Solutions

Application Platform

AI Builder



Power BI



Power Apps



Power Automate



Power Virtual Agents

Scenario-Based Services

Applied AI Services



Bot Service



Cognitive Search



Form Recognizer



Video Indexer



Metrics Advisor



Immersive Reader

Customizable AI Models

Cognitive Services



Vision



Speech



Language



Decision

OpenAI Service

ML Platform



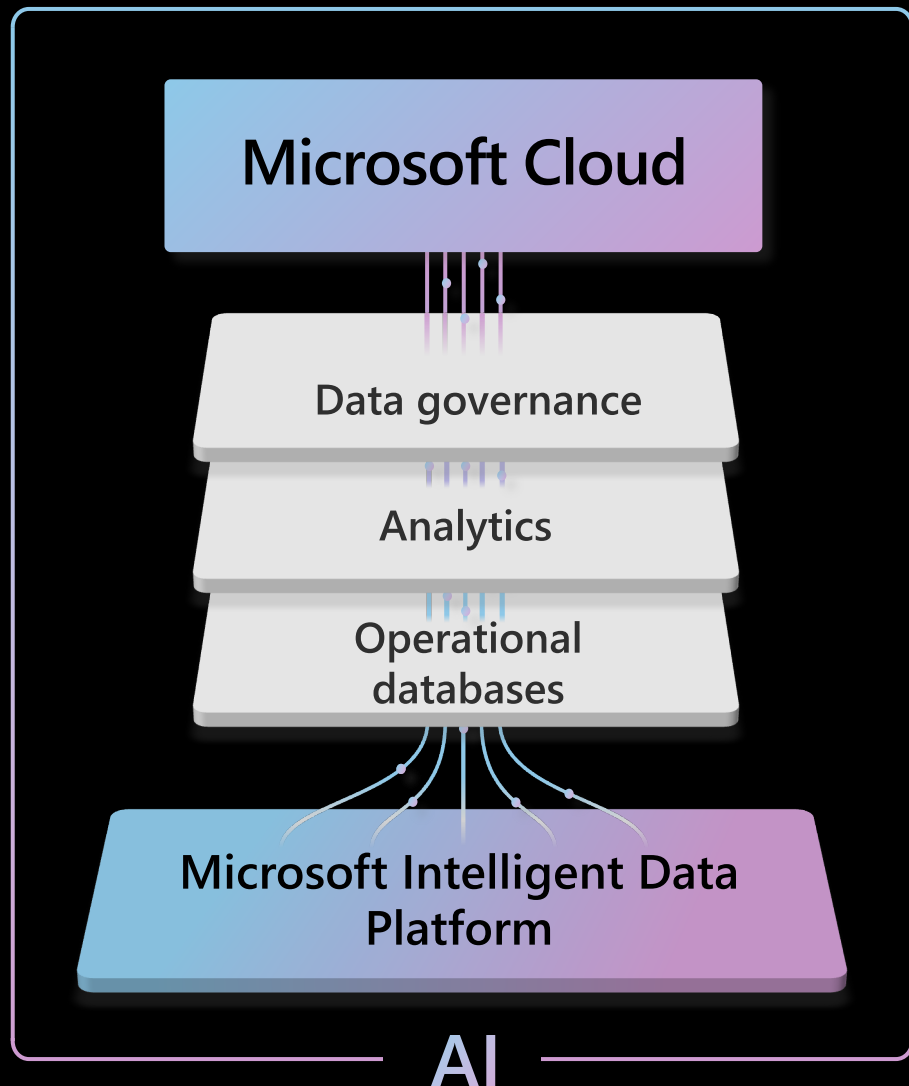
Azure Machine Learning




Business Users



Developers & Data Scientists



Deployed within your Azure subscription, secured by you, accessed only by you, and tied to your datasets and applications



Enterprise-grade security with role-based access control (RBAC) and authentication



Secure networking through private endpoints and VNETs

Generative AI

Content creation by API

Prompt

Write a tagline for an ice cream shop.

Response

We serve up smiles with every scoop!

Prompt

I'm having trouble getting my Xbox to turn on.

Response

There are a few things you can try to troubleshoot this issue

Prompt

Thanks! That worked. What games do you recommend for my 14-year-old?

Response

Here are a few games that you might consider: ...

Prompt

```
Table customers, columns =  
[CustomerId, FirstName,  
LastName, Company, Address,  
City, State, Country,  
PostalCode]
```

Create a SQL query for all customers in Texas named Jane
query =

Response

```
SELECT *  
FROM customers  
WHERE State = 'TX' AND  
FirstName = 'Jane'
```

Prompt

A ball of fire with vibrant colors to show the speed of innovation at our media and entertainment company

Response



Azure OpenAI | Concepts

Prompt and Completion

A short piece of text used to initiate an AI language model to generate a longer piece of text.

In computer science, the time complexity is the computational complexity that describes the amount of computer time it takes to run an algorithm. Time complexity is commonly estimated by counting the number of elementary operations performed by the algorithm, supposing that each elementary operation takes a fixed amount of time to perform. Thus, the amount of time taken and the number of elementary operations performed by the algorithm are taken to be related by a constant factor [...]

Summarize the previous text in a 2 sentences:

Prompt

Time complexity is a measure of the amount of time it takes for an algorithm to run, and is usually estimated by counting the number of elementary operations performed. It is usually expressed as a function of the size of the input, and the asymptotic behavior of the complexity is often focused on.

Completion

Best Practices

Be Specific

Be Descriptive

Double Down

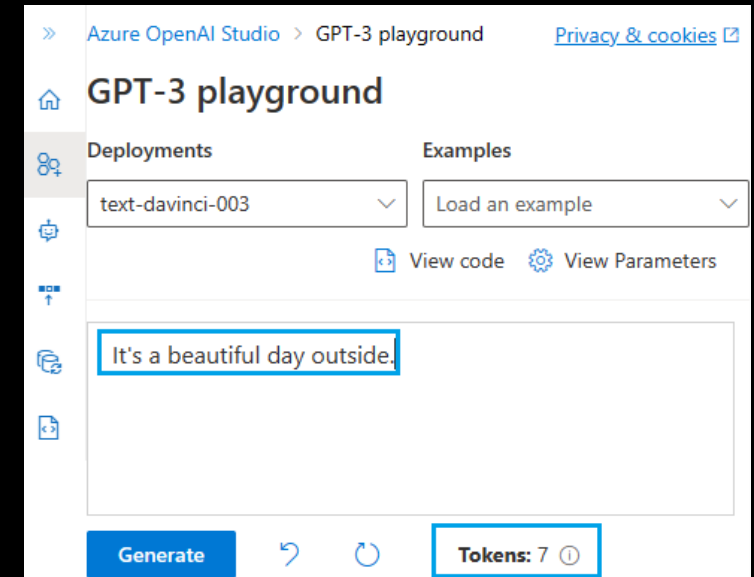
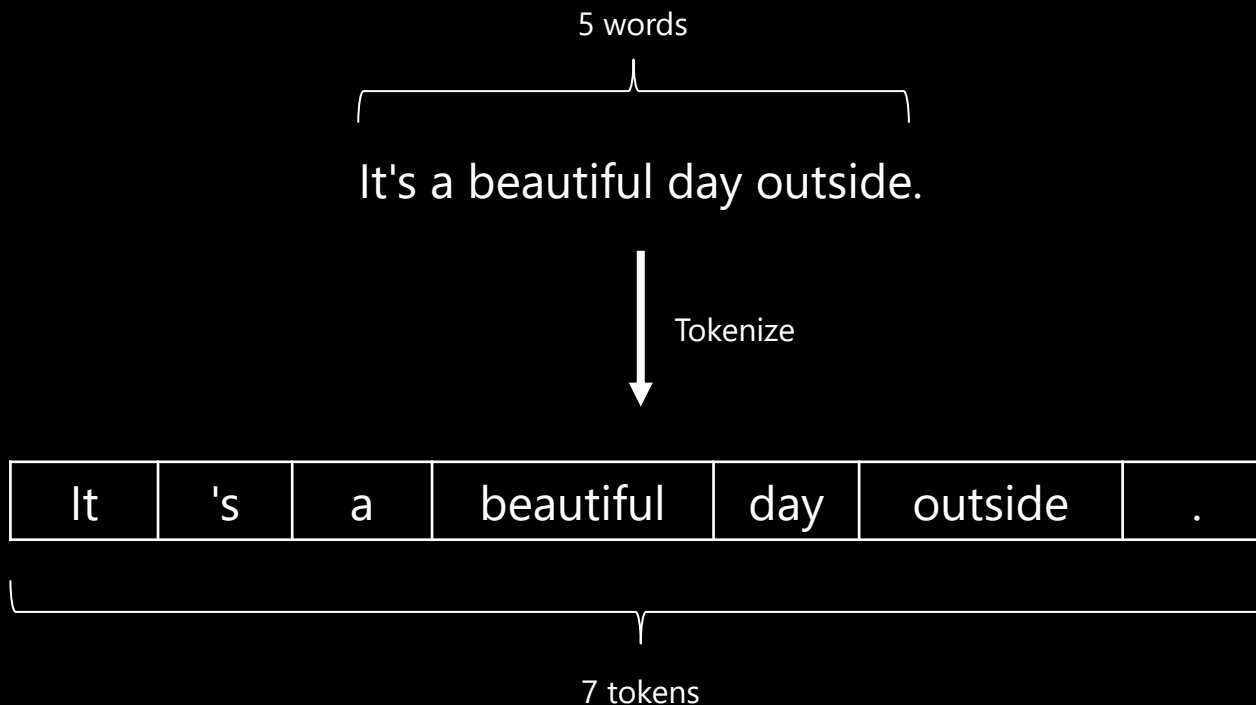
Order Matters

Alternative Path

Azure OpenAI | Concepts

Tokens

Smaller chunks of text created by splitting input text.



Azure OpenAI | Concepts

Embeddings

Representation of tokens as numerical vectors for machine learning algorithms to process.

It's a beautiful day outside.

5 words

↓ Tokenize

b'It'	b"'s"	b' a'	b' beautiful'	b' day'	b' outside'	b'.'
-------	-------	-------	---------------	---------	-------------	------

7 tokens

↓ Embeddings

[0.02132229320704937, -0.01934889890253544,, , 0.009578186087310314, -0.00010603990813251585]

1536 element vector

text-embedding-ada-002

1

Azure OpenAI | Concepts

Cosine Similarity

Metric used to measure similarity between two documents / text contents.

It's a beautiful day outside.

→ [.....]

Embeddings

The sun is shining brightly, and it is a perfect day for taking a stroll in the park.

→ [.....]

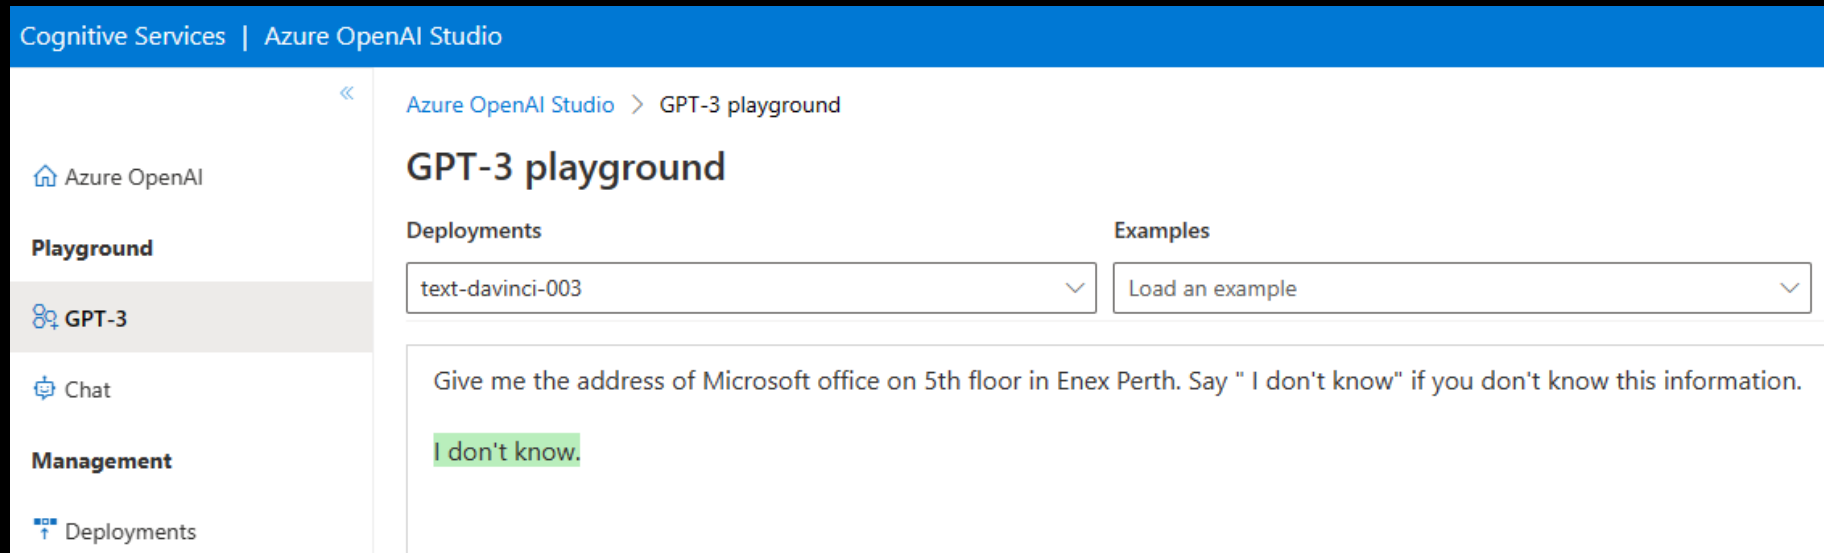
Cosine Similarity(v1, v2) = 0.92



Azure OpenAI | Concepts

Hallucinations

A confident response that does not seem to be justified by the model's training data
- Wikipedia



Azure OpenAI | Concepts

OpenAI Models

As of March 2023

Different pre-trained models, grouped by family and capability.

<i>More accurate, difficult problems</i>	<i>Understand and generate natural language</i>	<i>Generate code</i>	<i>Create embeddings</i>
<u>GPT-4 (Preview)</u>	<u>GPT-3</u>	<u>Codex</u>	<u>Embeddings</u>
gpt-4 gpt-4-32k	text-davinci-003 text-curie-001 text-babbage-001 text-ada-001 gpt-35-turbo (Preview)	code-davinci-002 code-cushman-001	text-similarity-ada-001 text-similarity-babbage-001 text-similarity-curie-001 text-similarity-davinci-001 ...
			text-search-ada-doc-001 text-search-ada-query-001 text-search-babbage-doc-001 ...
			code-search-ada-code-001 code-search-ada-text-001 ...

Azure OpenAI | Concepts

Model Parameters

User configurable settings that influence the model output.

Temperature, Top probabilities

Regulate the level of determinism in the model's response generation process.

Max response

No. of tokens to return as model response.

Frequency penalty

Control the probability of generating identical text in the model's response.

Presence penalty

Control token repetition to promote new topic introduction in responses.

Best of

Produce several responses and select the one with the highest overall probability of its tokens for display.

Azure OpenAI | Concepts

Zero-shot, Few Shot Learning

methods of using a model without any fine-tuning or additional training.

Zero-shot

Convert movie titles into emoji:

Spider-Man: Into the Spider-Verse

Result

Spider-Man: 🕷️ 🧔

Few-shot

Convert movie titles into emoji.

Back to the Future: 🧔 🧔 🚗 ⌚

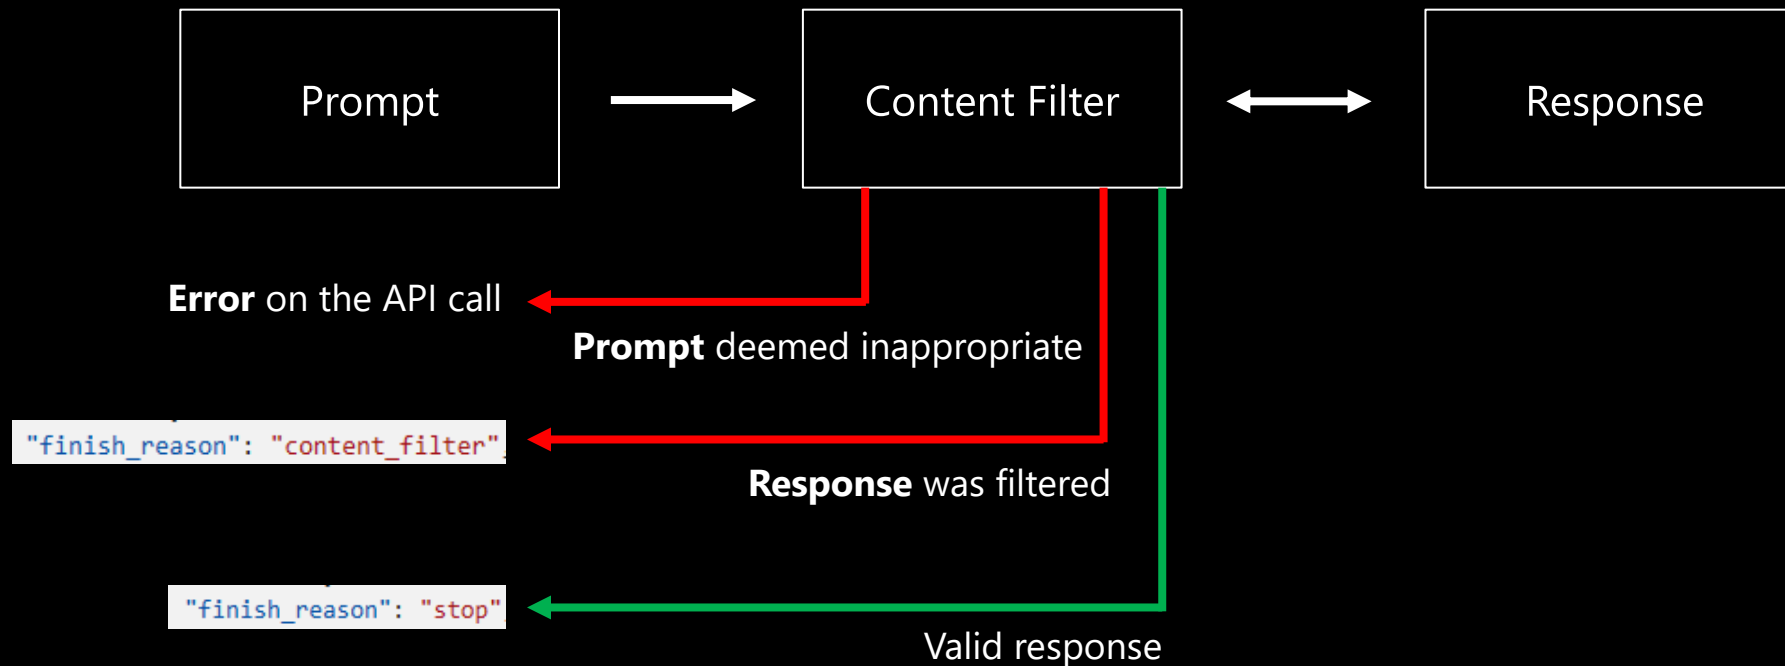
Batman: 🧔 🦇

Transformers: 🚗 🤖

Azure OpenAI | Concepts

Content filtering

Task of detecting misuse by analyzing prompt and response through classification models.

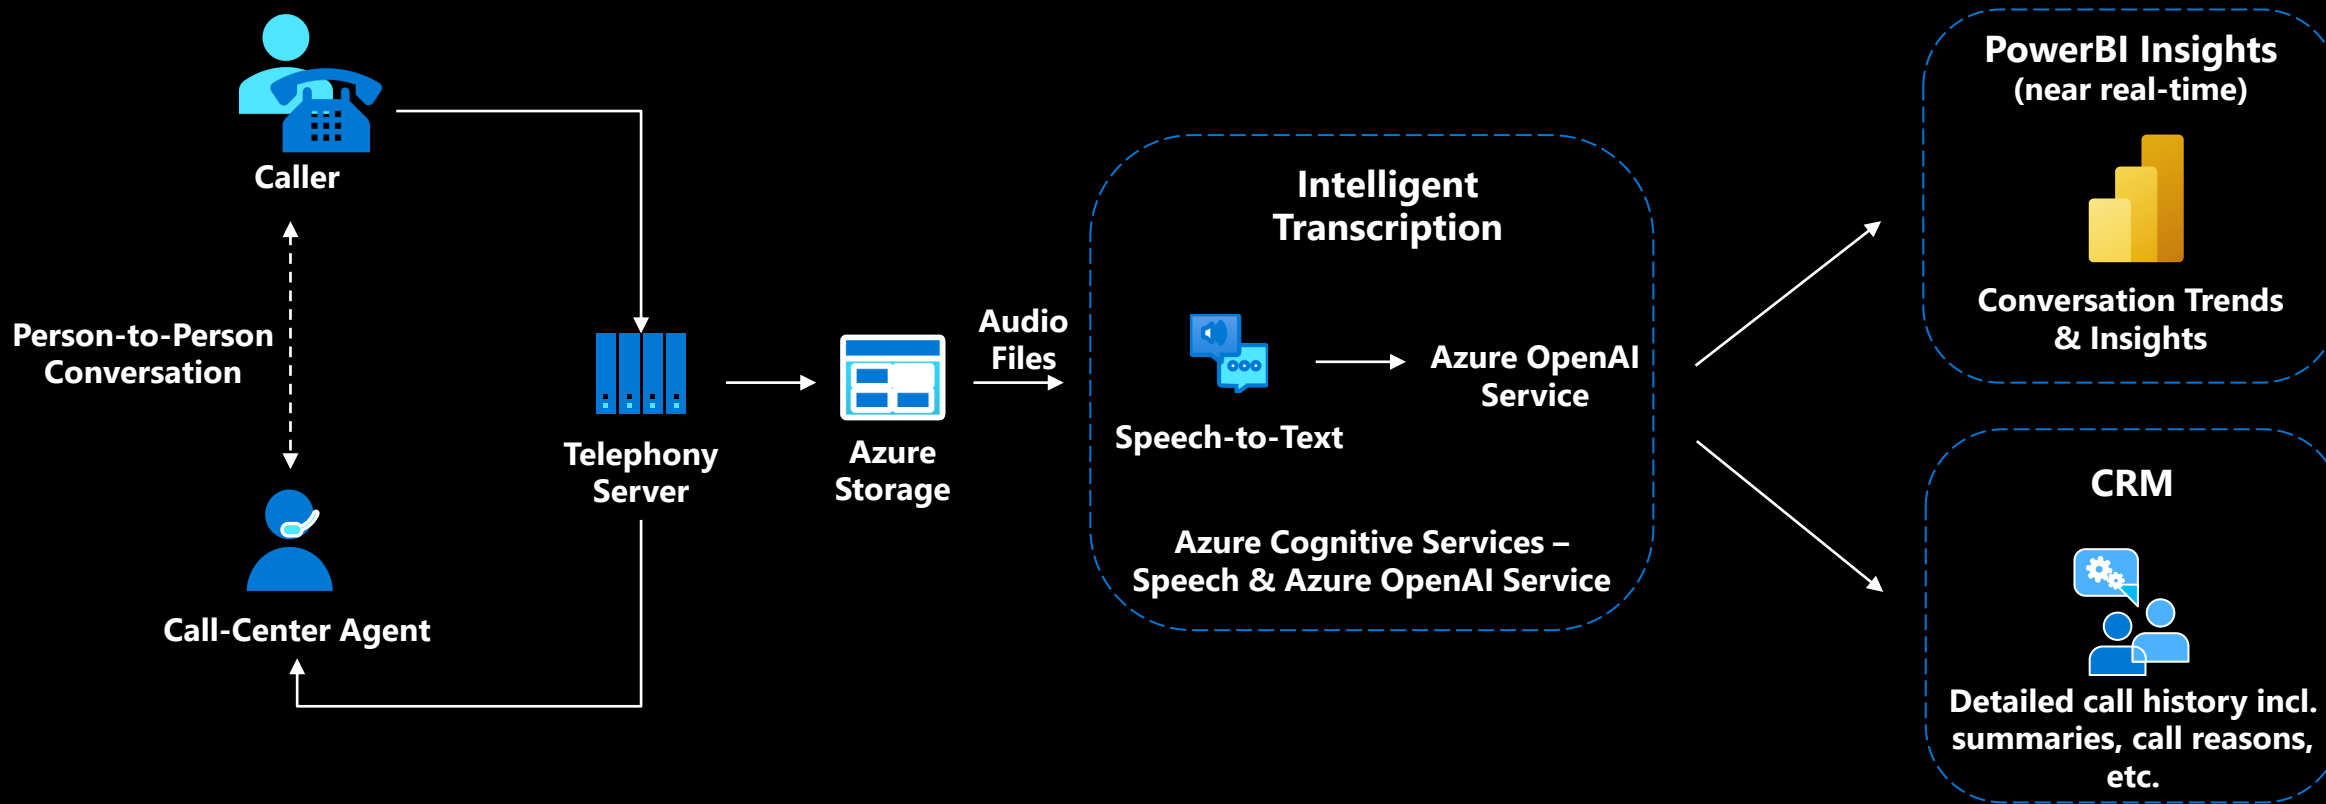




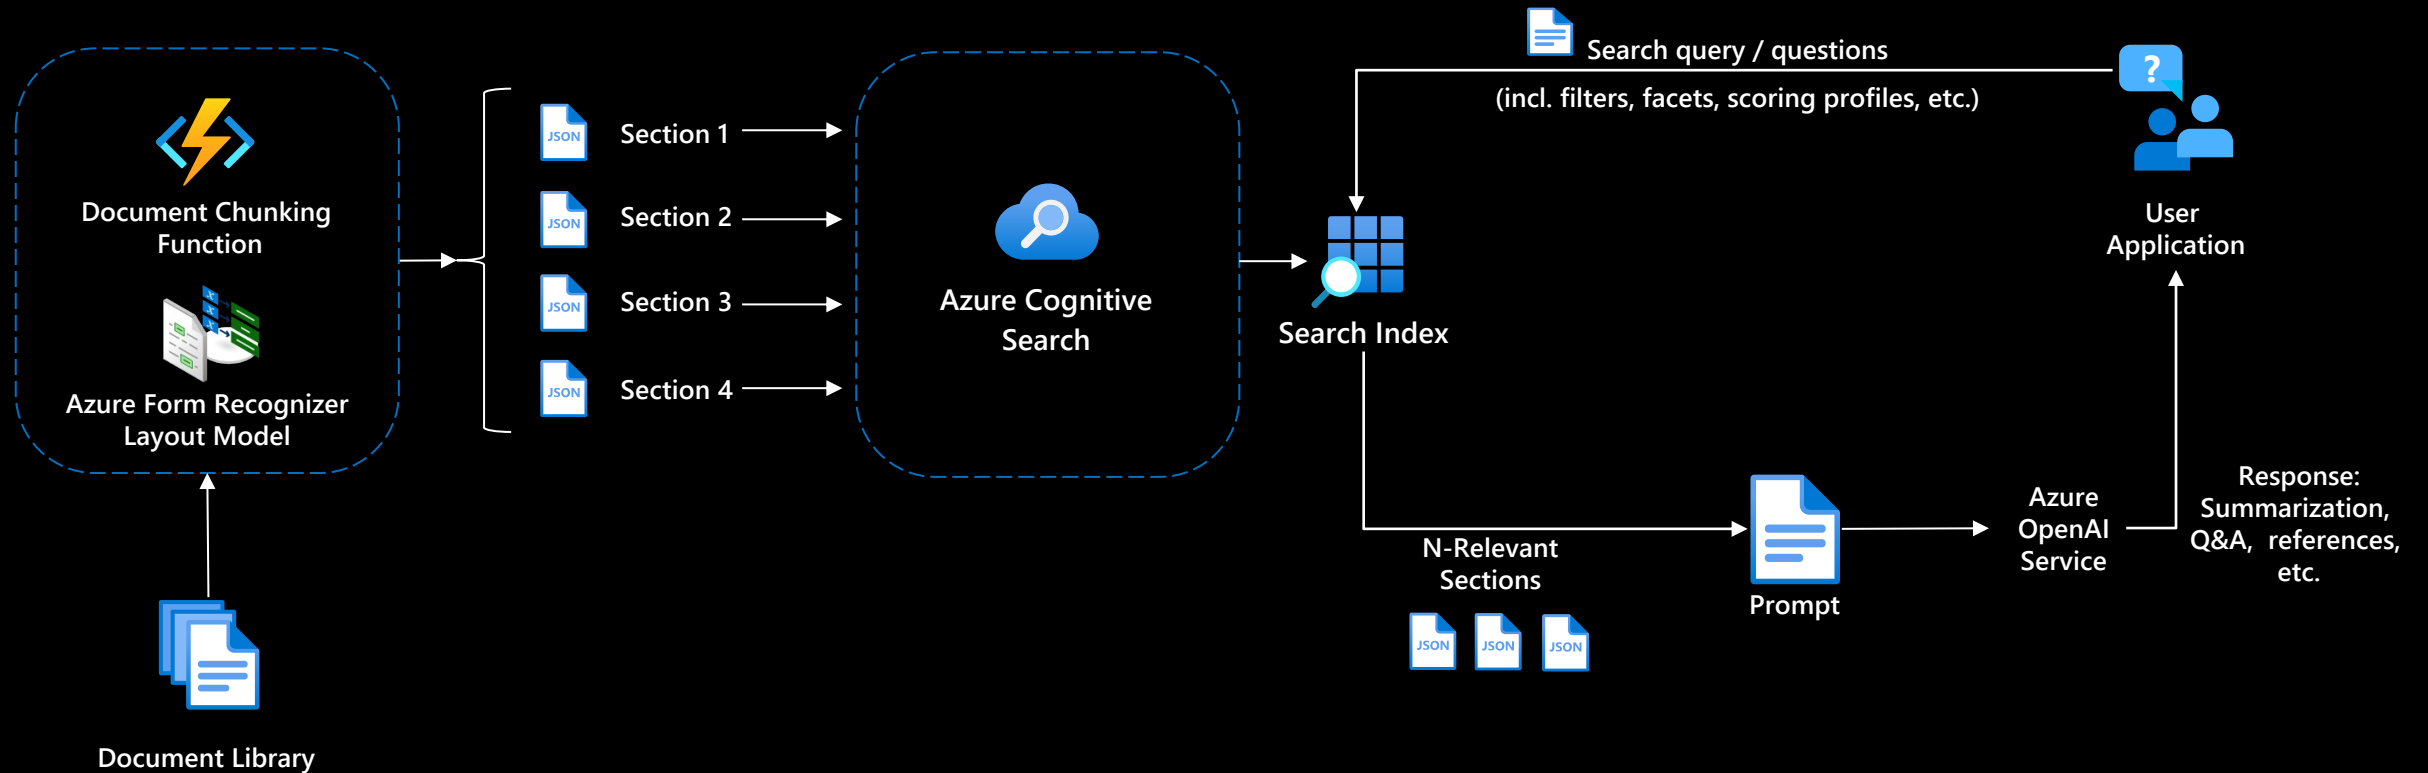
Demo – Azure OpenAI

Contact Center Analytics using Speech API & Azure OpenAI Service

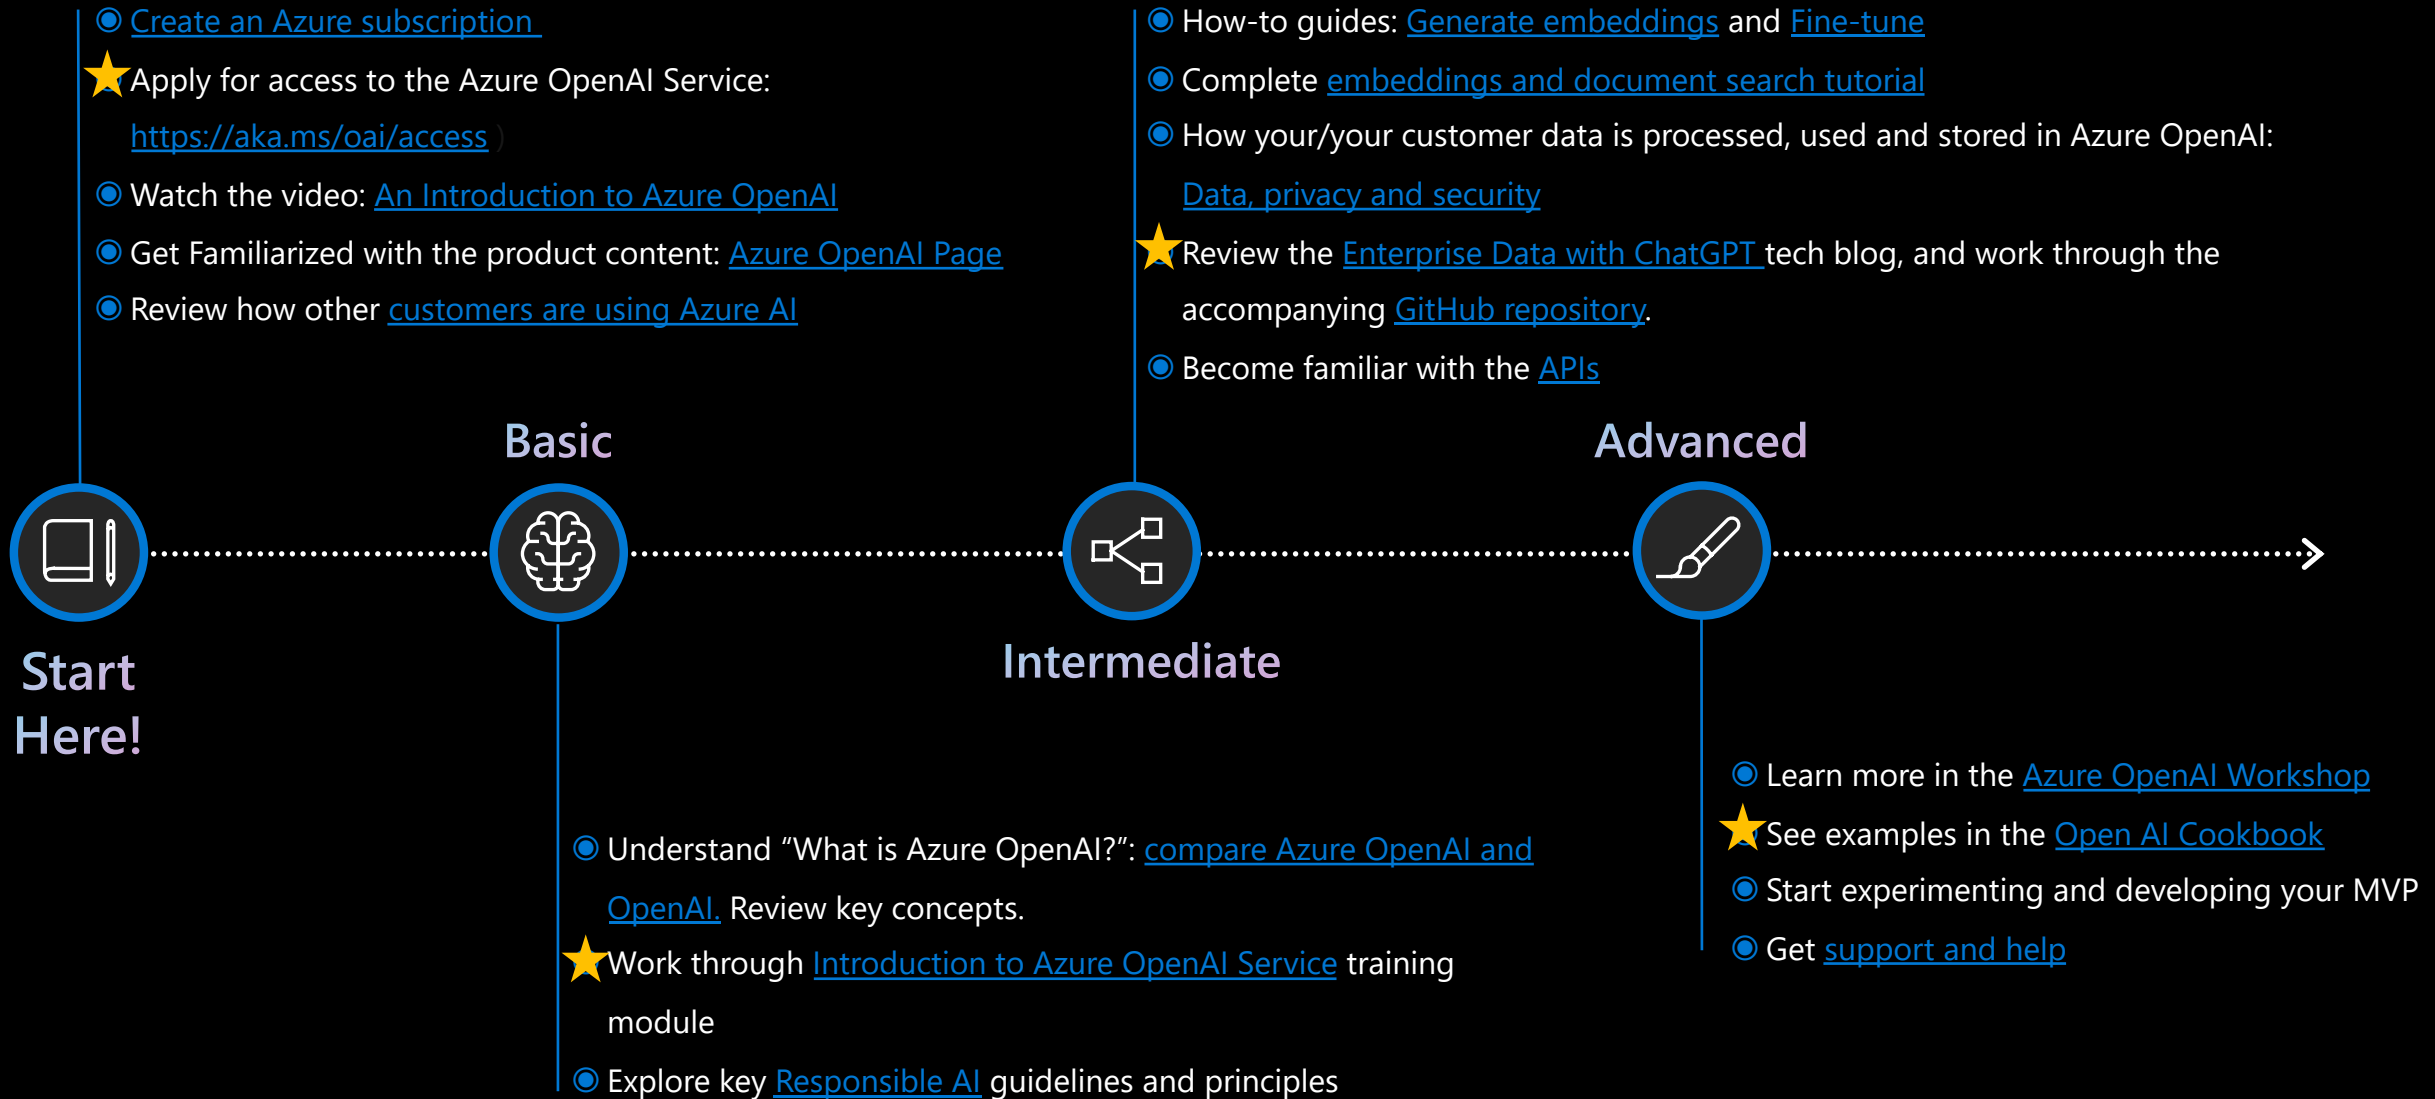
Extract rich insights from call transcripts



AI-Powered Q&A over Enterprise Data Sources



Azure OpenAI Service Learning Guide





In Seattle: May 23–25, 2023 | Online: May 23–24, 2023 PDT

Join us in Seattle for Microsoft Build!


President & Co-Founder Greg Brockman from OpenAI will be LIVE at Microsoft Build! Be there to explore how Microsoft and OpenAI's full-stack AI platform empowers developers and get ready to make AI work for you.

Register now

Showing 18 speakers



New content available!
Reload page



<https://build.microsoft.com/>



Q&A

Thank you!

Connect



in /sergiozenatti